

CS 188: Artificial Intelligence

Markov Decision Processes



Instructor: Stuart Russell and Dawn Song

University of California, Berkeley

Recap: Markov Decision Process (MDP)

- What is a Markov Decision Process?



Andrey Markov
(1856-1922)

Recap: Markov Decision Process (MDP)

- What is a Markov Decision Process?
 - State transition model is markov
 - Utility function is additive discounted rewards
- An MDP is defined by:
 - A set of states $s \in S$
 - A set of actions $a \in A$
 - A transition model $T(s, a, s')$
 - Probability that a from s leads to s' , i.e., $P(s' | s, a)$
 - A reward function $R(s, a, s')$ for each transition
 - A start state
 - Possibly a terminal state (or absorbing state)
 - Utility function which is additive discounted rewards

$$U_h([s_0, a_0, s_1, a_1, s_2, \dots]) = R(s_0, a_0, s_1) + \gamma R(s_1, a_1, s_2) + \gamma^2 R(s_2, a_2, s_3) + \dots$$


where $\gamma \in [0,1]$ is the **discount factor**



Andrey Markov
(1856-1922)


Recap: Policies

- What is a policy?

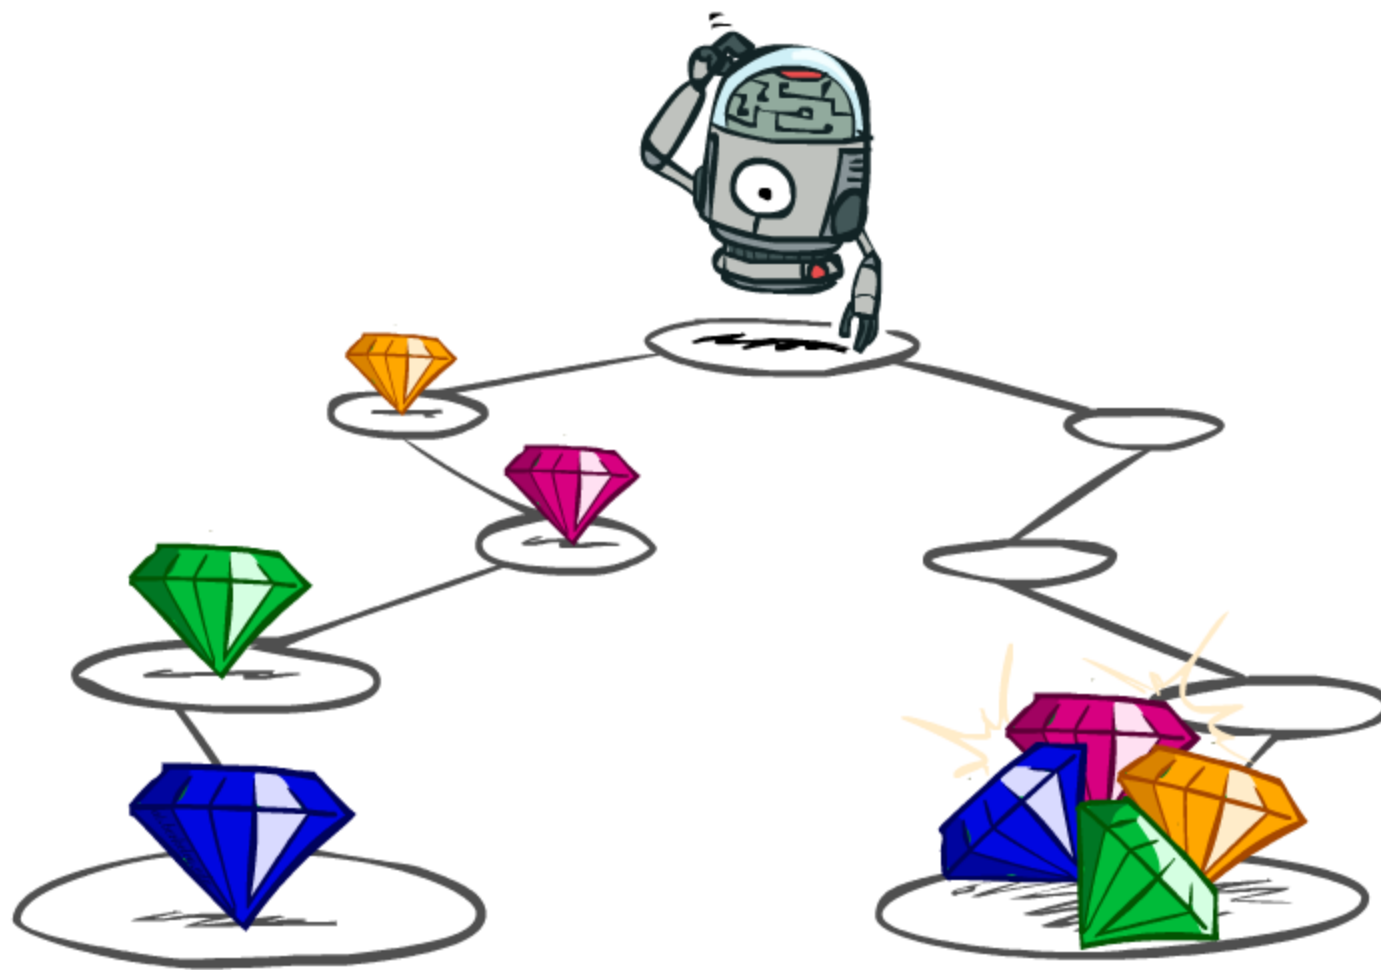


Recap: Policies

- A policy π gives an action for each state, $\pi: S \rightarrow A$
- For MDPs, we want an optimal **policy** $\pi^*: S \rightarrow A$
 - An optimal policy maximizes expected utility



Recap: Utilities of Sequences



Recap: Stationary Preferences

- Theorem: if we assume **stationary preferences**:


$$[s_0, a_0, s_1, a_1, s_2, \dots] > [s'_0, a'_0, s'_1, a'_1, s'_2, \dots] , \quad s_0 = s'_0, \quad a_0 = a'_0, \quad \text{and} \quad s_1 = s'_1 \\ \Leftrightarrow [s_1, a_1, s_2, \dots] > [s'_1, a'_1, s'_2, \dots]$$

then there is only one way to define utilities:

- **Additive discounted utility**:

$$U_h([s_0, a_0, s_1, a_1, s_2, \dots]) = R(s_0, a_0, s_1) + \gamma R(s_1, a_1, s_2) + \gamma^2 R(s_2, a_2, s_3) + \dots$$

where $\gamma \in [0,1]$ is the **discount factor**



Discounting



Worth r now



Worth γr next step



Worth $\gamma^2 r$ in two steps

- Discounting conveniently solves the problem of infinite reward streams!
 - Geometric series: $1 + \gamma + \gamma^2 + \dots = 1/(1 - \gamma)$
 - Assume rewards bounded by $\pm R_{\max}$
 - Then $r_0 + \gamma r_1 + \gamma^2 r_2 + \dots$ is bounded by $\pm R_{\max}/(1 - \gamma)$
- (Another solution: environment contains a **terminal state**; **and** agent reaches it with probability 1)

Quiz: Discounting

- Given:

10				1
a	b	c	d	e

- Actions: East, West, and Exit (only available in exit states a, e)
- Transitions: deterministic

- Quiz 1: For $\gamma = 1$, what is the optimal policy?

10				1
----	--	--	--	---

- Quiz 2: For $\gamma = 0.1$, what is the optimal policy?


10				1
----	--	--	--	---

- Quiz 3: For which γ are West and East equally good when in state d?

The utility of a policy


- Executing a policy π from any state s_0 generates a sequence
 $s_0, \pi(s_0), s_1, \pi(s_1), s_2, \dots$
- This corresponds to a sequence of rewards
 $R(s_0, \pi(s_0), s_1), R(s_1, \pi(s_1), s_2), \dots$
- This reward sequence happens with probability
 $P(s_1 \mid s_0, \pi(s_0)) \times P(s_2 \mid s_1, \pi(s_1)) \times \dots$
- The value (expected utility) of π in s_0 is written $U^\pi(s_0)$
 - It's the sum over all possible state sequences of
(discounted sum of rewards) \times (probability of state sequence)

$$U^\pi(s_0) = E \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(s_t), s_{t+1}) \right]$$




Optimal Quantities

- The optimal policy:
 $\pi^*(s)$ = optimal action from state s
Gives highest $U^\pi(s)$ for any π
- The value (utility) of a state s :
 $U^*(s) = U^{\pi^*}(s)$ = expected utility starting in s and acting optimally
- The value (utility) of a q-state (s,a) :
 $Q^*(s,a)$ = expected utility of taking action a in state s and (thereafter) acting optimally
 $U^*(s) \text{ ? } Q^*(s,a)$



Optimal Quantities

- The optimal policy:
 $\pi^*(s)$ = optimal action from state s
Gives highest $U^\pi(s)$ for any π
- The value (utility) of a state s :
 $U^*(s) = U^{\pi^*}(s)$ = expected utility starting in s and acting optimally
- The value (utility) of a q-state (s,a) :
 $Q^*(s,a)$ = expected utility of taking action a in state s and (thereafter) acting optimally
 $U^*(s) = \max_a Q^*(s,a)$



Bellman equations (Shapley, 1953)

- The value/utility of a state is
 - The expected reward for the next transition plus the discounted value/utility of the next state, assuming the agent chooses the optimal action
- Hence we have a recursive definition of value (Bellman equation):

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U(s')]$$


- Similarly, Bellman equation for Q-functions

$$\begin{aligned} Q(s, a) &= \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U(s')] \\ &= \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma \max_{a'} Q(s', a')] \end{aligned}$$

Solving MDPs

- Value iteration
- Policy iteration

Value Iteration




Value Iteration

- Start with (say) $U_0(s) = 0$ and some termination parameter ϵ
- Repeat until convergence (i.e., until all updates smaller than ϵ)
 - Do a **Bellman update** (essentially one ply of expectimax) from each state:

$$U_{k+1}(s) \leftarrow \max_a \sum_{s'} P(s' | a, s) [R(s, a, s') + \gamma U_k(s')]$$

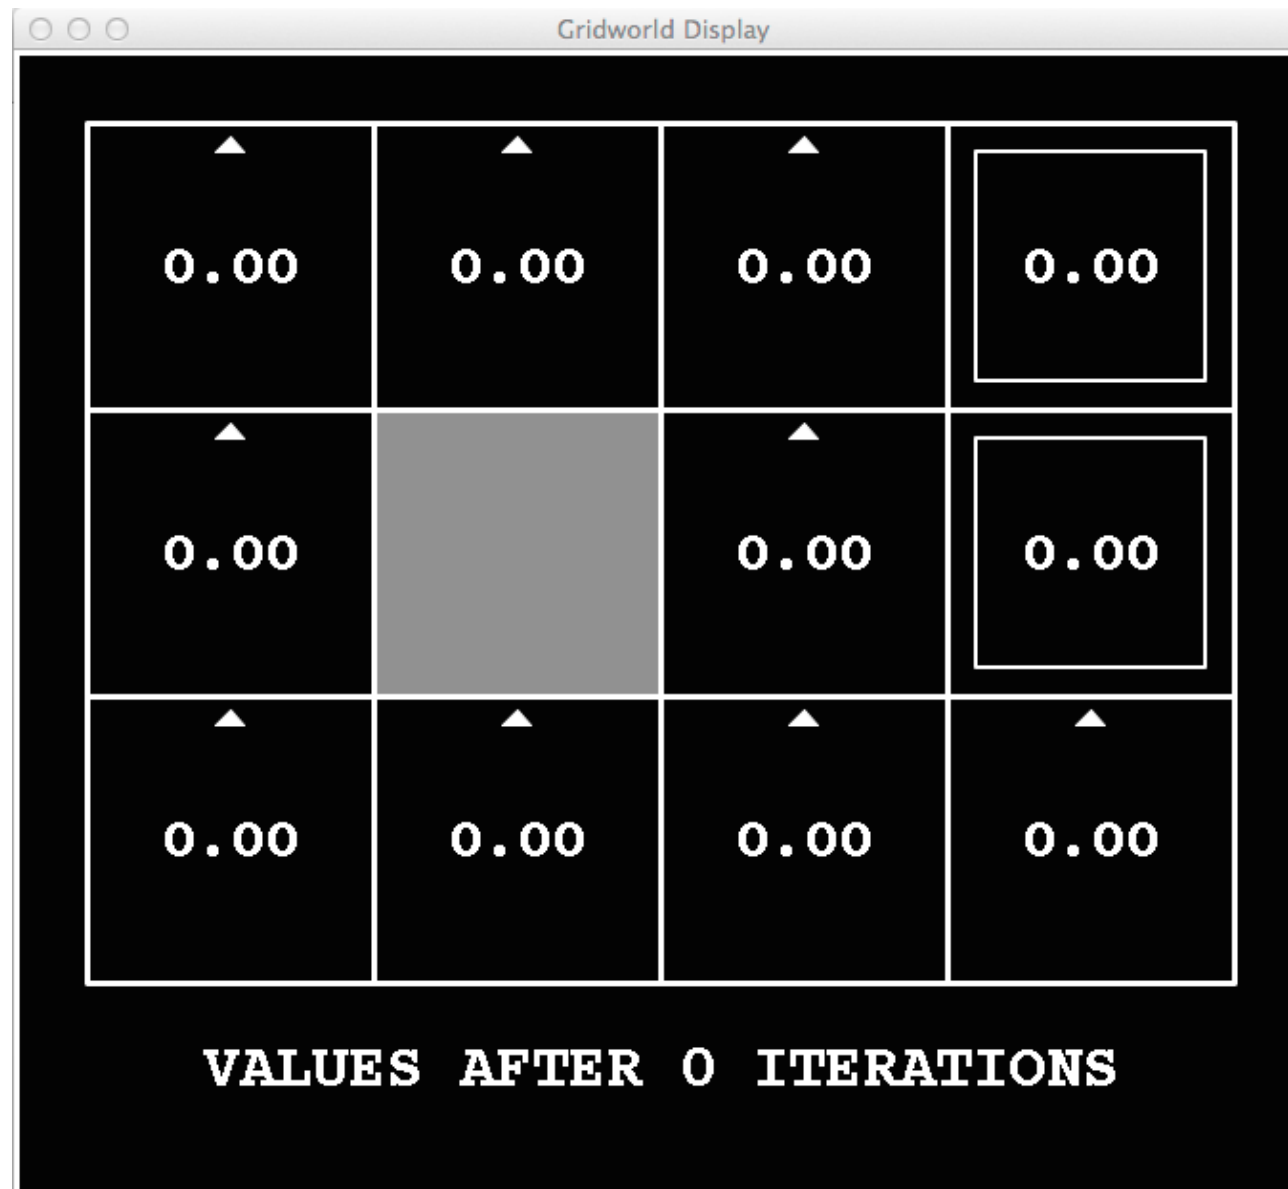
$$U \leftarrow BU$$

- Theorem: will converge to unique optimal values
- Runtime per iteration?
 - Complexity of each iteration: $O(S^2A)$



k=0


$$U_{i+1}(s) \leftarrow \max_{a \in A(s)} \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U_i(s')]$$



Noise = 0.2
Discount = 0.9
Living reward = 0

k=1

$$U_{i+1}(s) \leftarrow \max_{a \in A(s)} \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U_i(s')]$$



Noise = 0.2
Discount = 0.9
Living reward = 0

k=2



Noise = 0.2
Discount = 0.9
Living reward = 0

$k=3$



Noise = 0.2
Discount = 0.9
Living reward = 0

$k=4$



Noise = 0.2
Discount = 0.9
Living reward = 0

k=5



Noise = 0.2
Discount = 0.9
Living reward = 0

k=6



Noise = 0.2
Discount = 0.9
Living reward = 0

$k=7$



Noise = 0.2
Discount = 0.9
Living reward = 0

k=8



Noise = 0.2
Discount = 0.9
Living reward = 0

k=9



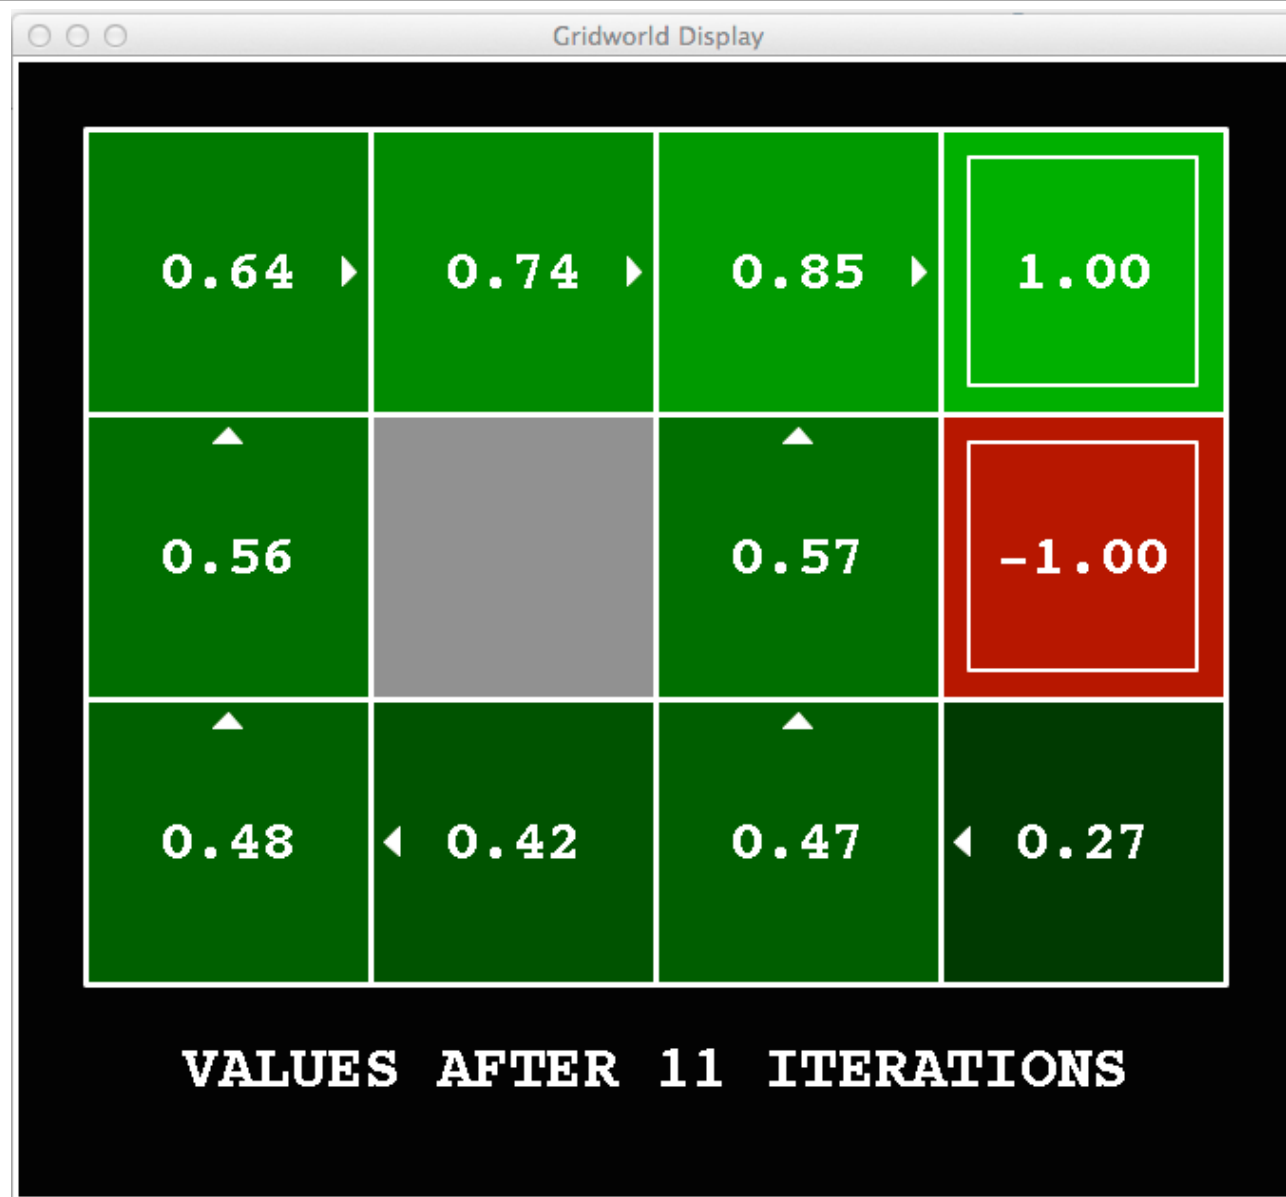
Noise = 0.2
Discount = 0.9
Living reward = 0

k=10




Noise = 0.2
Discount = 0.9
Living reward = 0

k=11



Noise = 0.2
Discount = 0.9
Living reward = 0

k=12



Noise = 0.2
Discount = 0.9
Living reward = 0

k=100



Noise = 0.2
Discount = 0.9
Living reward = 0